GNN-Augmented BP Architectures for QLDPC Decoding

Jothiradithya Sai Konduru[1], Anish Chedalla[1], and Nithin Raveendran[2]

[1]Paradise Valley High School, Phoenix, AZ, USA
[2]Department of Electrical and Computer Engineering, University of Arizona, Tucson, AZ, USA

Bivariate bicycle (BB) quantum low-density parity-check (QLDPC) codes combine finite encoding rate with low-weight checks, but their loopy Tanner graphs remain challenging for iterative decoding. ~~Learned message-passing can improve belief propagation (BP), although the measured gain depends strongly on the BP configuration used for training and evaluation. We study this issue using~~ We study a Tanner-graph neural network (TannerGNN) that ~~produces~~ adds qubit-wise log-likelihood corrections ~~for BPunder a biased CSS component-noise model. On the $[\![72, 12, 6]\!]$ BB code, a model trained~~ to belief propagation (BP) under biased CSS noise. The network is trained with a per-qubit focal loss on the sampled error, and its checkpoint is selected with 10-iteration flooding BP. On the $[\![72, 12, 6]\!]$ BB code at $p = 0.04$ and $\eta = 20$, it reduces the logical error rate (LER) of that decoder by $12.5\% \pm 0.6\%$ on 20,000 held-out syndromes ~~at $p = 0.04$ and $\eta = 20$~~ (five training seeds; paired McNemar $p < 10^{-36}$ for every seed). ~~This~~ The gain does not transfer to a stronger reference decoder, serial min-sum BP-OSD~~decoder: the learned~~ : the correction never improves it and, after Holm correction, is significantly worse ~~at $p \geq 0.03$. A paired failure-set analysis shows that the failure sets of the two decoders are nested and that most shots the GNN repairs are already decoded by~~ for 11 of the 15 seed and error-rate pairs with $p \geq 0.03$. Most BP-OSD failures are also BP failures, the GNN mainly repairs shots on which BP converges slowly and which BP-OSD already decodes, and weight-three errors that cover half of a stabilizer defeat BP with or without the GNN but never BP-OSD. ~~Independent baseline experiments further show that schedule, BP~~ Baseline experiments show that BP schedule, update rule, iteration count ~~, and OSD configuration can change~~ and OSD settings change the LER by orders of magnitude~~. We also establish that unclipped min-sum BP is invariant, in its hard decisions, to a uniform positive rescaling of all channel LLRs within a CSS component. Thus, changing only the magnitude of a uniform prior cannot make the underlying~~ , and we prove that uniformly rescaling the channel LLRs cannot change the hard decisions of unclipped min-sum BP~~adaptive, although a learned network may still use global context through other pathways. Finally, synthetic experiments with heterogeneous~~ . Revealing the true per-qubit error rates~~show a large gap between mean-prior and oracle-prior decoding, motivating spatially resolved calibration side information rather than a single global noise scalar. The experiments show that an improvement over one BP baseline does not by itself establish an improvement over a stronger decoder, and that spatial calibration information is qualitatively different from a single global noiseestimate~~, by contrast, reduces the BP-OSD LER by up to 88% under synthetic heterogeneous noise and by 55% under circuit-level noise.

Nithin Raveendran: nithin@arizona.edu

1 Introduction

Bivariate bicycle (BB) codes provide a concrete finite-length setting in which QLDPC codes combine non-vanishing encoding rate, weight-six stabilizer checks, and syndrome-extraction circuits compatible with low-degree connectivity [1]. IBM's current fault-tolerant architecture is explic-

itly based on BB codes, making the finite-length decoding behavior of this family practically relevant as well as theoretically interesting [2][2, 3]. The canonical examples include the $[\![72, 12, 6]\!]$, $[\![144, 12, 12]\!]$, and $[\![288, 12, 18]\!]$ codes studied here.

The decoding bottleneck is not the absence of efficient iterative algorithms, but their finite-length failure modes. BP is attractive because its message updates are local and parallelizable, yet short cycles and degeneracy in QLDPC Tanner graphs create correlated messages and non-convergent or incorrect fixed points [4]. OSD post-processing can recover many BP failures, while recent learned decoders use graph neural networks (GNNs) either to replace BP or to intervene between BP stages [5, 6]. These approaches make the choice of baseline unusually consequential: a neural correction ~~trained against~~ developed for one BP schedule or stopping rule need not address the residual failures of a stronger decoder.

A second issue is adaptivity. Hardware noise is neither spatially nor temporally uniform, and calibration information can therefore be useful to a decoder. However, three distinct sources of information should not be conflated: a global estimate of the physical error rate, spatially resolved per-qubit calibration data, and syndrome history. A decoder may respond differently to each. In particular, changing the magnitude of a uniform channel prior has a restricted effect on min-sum BP, whereas per-qubit priors can change the relative reliability ordering that drives both BP and OSD.

We study a GNN-augmented BP pipeline with two questions in view: *what gain does the learned correction provide relative to the BP decoder ~~used in training~~it is developed with, and what additional information is needed for noise adaptation?* We compare decoder configurations, training and evaluation baselines, and different forms of calibration information to isolate the source of the observed performance changes.

1.1 Related work

Quantum Tanner graphs combine short cycles with degeneracy, and both cause non-convergence and finite-length error floors in BP [4, 7, 8]. Trapping-set analysis makes these failures structural rather than mysterious: specific subgraphs, including symmetric configurations created by stabilizers, can trap iterative decoders [9], and degeneracy itself can be exploited inside BP [10].

Several families of methods repair such failures, with different latency and complexity trade-offs. OSD post-processing [4, 11, 12] is the most widely used reference. Alternatives include stabilizer inactivation [13], check-agnosia [14], localized statistics decoding [15], syndrome-based linear programming [16], collective bit flipping [17], and trapping-set-aware iterative decoders of linear complexity [18]. For BB codes in particular, recent work targets real-time decoding: Relay-BP [19], graph rewiring with ensembles of normalized min-sum decoders [20], diversity and ensemble methods [21], ambiguity clustering [22], and hardware implementations of BP-OSD [23].

Learned decoders either modify BP messages or take over part of the inference. Neural BP with a degeneracy-aware loss [24], GNN stages interleaved with BP [5], learned message passing [6], circuit-level neural decoders for BB codes [25], and calibration-conditioned decoders [26] report both strong gains and limits in scaling and transfer. In this landscape serial min-sum BP-OSD is a useful, fully specified reference decoder, not the state of the art, and we use it only in that role. Our question is narrower: what a learned LLR correction gains relative to the BP configuration it is developed with, whether that gain survives a stronger reference decoder, and what side information adaptation actually needs.

1.2 Contributions

Our main findings are:

1. **A held-out GNN gain over ~~its training-aligned~~ the BP ~~baseline~~decoder it is developed with.** On $[\![72, 12, 6]\!]$ at $p = 0.04$ and $\eta = 20$, TannerGNN+flooding BP reduces LER by $12.5\% \pm 0.6\%$ on 20,000 held-out syndromes over five training seeds (McNemar $p < 10^{-36}$ for every seed). The training rates span $p \in [0.005, 0.045]$; at $p = 0.06$, outside that range, the reduction is $10.6\% \pm 0.7\%$. The whole gain is in the Z component; adding the correction to the X component as well only adds failures there (Sec. 5.3).

2. **Baseline configuration and alignment can dom-**

inate the apparent learned gain. Serial min-sum BP with OSD post-processing is substantially stronger than the flooding BP used ~~in GNN training~~to select the GNN checkpoint. The same learned correction ~~does not improve~~ never improves this stronger decoder and, after Holm correction, is significantly worse ~~from $p = 0.03$ on~~for 11 of the 15 seed and error-rate pairs with $p \geq 0.03$. A paired failure-set analysis shows that 93% of BP-OSD failures are also BP failures (61% the other way) and that most shots the GNN repairs are already decoded by BP-OSD. Separate experiments on $[\![288, 12, 18]\!]$ show order-of-magnitude sensitivity to BP schedule, update rule, iteration count, and OSD settings. The ~~training decoder and~~ decoder used for model selection and the evaluation decoder are therefore part of the experimental condition rather than interchangeable implementations of "BP" (Secs. 5.1 and 5.4).

3. **A recurring BP failure that the GNN does not repair.** When a weight-three Z error covers half of a weight-six stabilizer, the error and its stabilizer complement are equally likely and logically equivalent. BP, 100-iteration BP and GNN-BP fail on every such shot in our test sets and BP-OSD on none; the GNN's bitwise loss gives both halves the same target, so it cannot learn to break the tie (Secs. 3.2 and 5.4).

4. **Global prior scaling and spatial calibration carry different information.** We prove that hard decisions of unclipped min-sum BP are invariant to a uniform positive rescaling of the channel LLRs within a CSS component. This result limits adaptivity obtained only by changing a uniform BP prior; it does not rule out a learned network using a global noise scalar as context. In synthetic heterogeneous-noise experiments, providing the true per-qubit rates to BP-OSD yields large oracle improvements, motivating spatially resolved calibration inputs. Because code size and operating point co-vary in these experiments, we do not attribute the cross-code trend to distance alone (Secs. 5.2 and 5.5).

2 Background

2.1 CSS codes and bivariate bicycle construction

A CSS code [27, 28] is defined by matrices $\mathbf{H}_x, \mathbf{H}_z \in \mathbb{F}_2^{m \times n}$ with $\mathbf{H}_x \mathbf{H}_z^\top = \mathbf{0}$. The syndrome of an error (e_x, e_z) decomposes as $s_x = \mathbf{H}_x e_z$, $s_z = \mathbf{H}_z e_x$ over $\mathbb{F}_2$, so X- and Z-type errors are decoded independently. This CSS independence is structurally important for the results of Sec. 5.2.

Bivariate bicycle codes [1] are constructed over the group algebra $\mathbb{F}_2[x, y]/(x^\ell - 1, y^m - 1)$. Polynomials $A(x, y)$ and $B(x, y)$ yield circulant block matrices:

$$\mathbf{H}_x = [\mathbf{A} \mid \mathbf{B}], \qquad \mathbf{H}_z = [\mathbf{B}^\top \mid \mathbf{A}^\top]. \tag{1}$$

We study three codes (Table 1).

2.2 Belief propagation and its failure modes

Min-sum BP on a CSS Tanner graph iterates variable-to-check and check-to-variable messages for T steps. For a check c_j with measured syndrome bit s_j and a neighboring variable v_i,

$$\mu^{(t)}_{v_i \to c_j} = \Lambda_i + \sum_{j' \in \mathcal{N}(i) \setminus j} \mu^{(t-1)}_{c_{j'} \to v_i}, \tag{2}$$

$$\mu^{(t)}_{c_j \to v_i} = (-1)^{s_j} \prod_{i' \in \mathcal{N}(j) \setminus i} \operatorname{sgn}\left(\mu^{(t)}_{v_{i'} \to c_j}\right) \times \min_{i' \in \mathcal{N}(j) \setminus i} \left|\mu^{(t)}_{v_{i'} \to c_j}\right|. \tag{3}$$

~~The channel LLR for qubit i under Z-biased noise is $\Lambda_{z,i} = \ln((1 - p_z)/p_z)$, $p_z = p\eta/(\eta + 1)$.~~ CSS decoding separates the X and Z components, processing each with its respective parity-check matrix. Under the noise model of Sec. 2.3 the matched channel LLRs are $\Lambda_{z,i} = \ln((1 - p_z)/p_z)$ for Z errors and $\Lambda_{x,i} = \ln((1 - p_x)/p_x)$ for X errors; at $p = 0.04$ and $\eta = 20$, $p_z = 0.0381$, $p_x = 0.0019$ and $\Lambda_x/\Lambda_z = 6.26/3.23 \approx 1.94$. Each experiment below states which prior every decoder receives.

The Tanner graphs of the finite-length BB codes considered here contain many short cycles. Message reuse around these cycles creates statistical dependencies that are absent on a tree and can lead BP to oscillate, converge to a syndrome-inconsistent estimate, or converge

Table 1: Bivariate bicycle codes studied in this work, with generator polynomials from [1]; $\mathbf{H}_x$ and $\mathbf{H}_z$ follow Eq. (1) over $\mathbb{F}_2[x,y]/(x^\ell-1,\, y^m-1)$. The names "gross" and "two-gross" follow IBM's usage [3].

Code	Name	n	k	d	Rate	ℓ	m	$A(x,y)$	$B(x,y)$
$[\![72,12,6]\!]$	—	72	12	6	1/6	6	6	x^3+y+y^2	y^3+x+x^2
$[\![144,12,12]\!]$	gross	144	12	12	1/12	12	6	x^3+y+y^2	y^3+x+x^2
$[\![288,12,18]\!]$	two-gross	288	12	18	1/24	12	12	$x^3+y^2+y^7$	y^3+x+x^2

to the wrong logical class. We therefore treat short-cycle structure as a finite-length decoder mechanism to be measured, rather than attributing every cycle directly to the CSS orthogonality condition. Post-processing methods including ordered-statistics decoding (OSD) [11] recover many such failures but add $O(n^3)$ worst-case complexity. As shown in Sec. 5.1, the specific OSD *configuration*—schedule, base method, and iteration count—can change LER by more than two orders of magnitude on the same code. We therefore report these settings explicitly in the comparisons below.

2.3 Noise models

We use a Z-biased CSS component-noise model with $\eta = 20$:

$$p_z = \frac{p\eta}{\eta+1}, \qquad p_x = \frac{p}{\eta+1}. \tag{4}$$

Each data qubit independently suffers a Z flip with probability p_z and an X flip with probability p_x. No separate Y channel is sampled: a qubit carries a Y error only if it receives both flips, which happens with probability $p_x p_z \approx p^2/22$, so this is not depolarizing noise. Syndromes are $s_x = \mathbf{H}_x e_z$ and $s_z = \mathbf{H}_z e_x$, and the two components are decoded separately. ~~For temporal drift experiments we apply three modulation models to the base rate p: sinusoidal ($p(t) = p + A\sin(2\pi t/T)$), Ornstein–Uhlenbeck, and random telegraph noise (RTN), all motivated by TLS physics [29, 30]. The decoder receives no direct access to instantaneous noise parameters in any~~ Two kinds of temporal variation appear, both motivated by fluctuating two-level systems [29, 30]: the GNN training data modulate the global rate sinusoidally, $p(t) = p_0 + A\sin(2\pi t/T)$ (Sec. 3.2), and the drift experiment of Sec. 5.6 lets the rate of each qubit follow its own Ornstein–Uhlenbeck process. Which noise parameters a decoder sees is stated with each experiment.

3 The AdaptiveBP Framework

This section describes the model evaluated in Sec. 5. Two further variants, FiLM conditioning and an interleaved GNN-BP pipeline, are described in Appendix A; they are not evaluated.

3.1 TannerGNN

TannerGNN operates on the unified CSS Tanner graph, with edges distinguished by type $\tau \in \{X\text{-check}, Z\text{-check}\}$. Each node receives a 4-dimensional feature vector whose first entry is a value and whose last three entries one-hot encode the node type: a variable node v_i receives $\mathbf{x}_i = [\ell, 1, 0, 0]$, ~~where $\ell = \ln((1-p)/p)$ is the prior LLR,~~ and an X-check (Z-check) node c_j receives $[s_j, 0, 1, 0]$ ($[s_j, 0, 0, 1]$), where s_j is its syndrome bit. The prior LLR $\ell = \ln((1-p)/p)$ is computed from the total physical error rate p of the shot, not from p_z or p_x: for the drifting training streams it is the instantaneous rate $p(t)$, and for the test sets the fixed p. The same ℓ is the channel prior of both CSS components in the flooding-BP decoders of Sec. 5. Its magnitude does not change min-sum decisions (Sec. 5.2), but it sets the scale against which the correction Δ_i acts.

An input projection maps features to hidden dimension d_h ($d_h = 32$ for the model of Sec. 5.3) via $\mathbf{h}_i^{(0)} = \mathrm{ReLU}(\mathrm{LayerNorm}(\mathbf{W}_{\mathrm{proj}}\mathbf{x}_i))$. Each of $L = 3$ message-passing layers applies separate edge-type MLPs, optional attention weighting, and a GRU node update:

$$\mathbf{m}_i^{(\ell)} = \sum_{j \in \mathcal{N}(i)} \mathrm{MLP}^{(\ell)}_{e,\tau(i,j)}([\mathbf{h}_i^{(\ell)}, \mathbf{h}_j^{(\ell)}]), \tag{5}$$

$$\mathbf{h}_i^{(\ell+1)} = \mathrm{GRU}(\mathbf{h}_i^{(\ell)},\, \mathbf{m}_i^{(\ell)}). \tag{6}$$

A readout MLP maps each variable node's final hidden state to a per-qubit LLR correction: $\Delta_i = \mathrm{MLP}_{\mathrm{read}}(\mathbf{h}_i^{(L)})$.

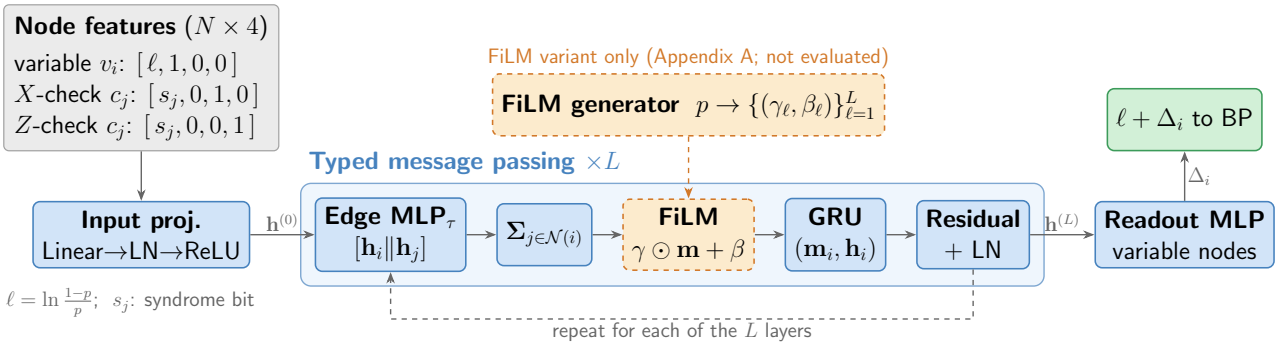


Figure 1: TannerGNN architecture. Variable and check nodes exchange typed messages over $L = 3$ layers via edge-type-specific MLPs; a GRU update with a residual connection and layer normalisation follows at each node, and a readout MLP on the variable nodes gives the LLR corrections. In the FiLM variant (dashed; Appendix A), FiLM modulation scales the aggregated messages at each layer; the model of Sec. 5.3 does not use FiLM.

~~The current implementation applies the same scalar Δ_i to both $A_{z,i}$ and $A_{x,i}$, and computes GNN input features from a single prior LLR ℓ. Under $\eta = 20$, $A_{z,i}$ and $A_{x,i}$ differ substantially (factor ≈ 1.7 in LLR magnitude at $p = 0.04$), so a single correction cannot simultaneously be optimal for both error types. Separate Z- and X-readout heads provide a direct way to remove this coupling and should be tested under the same biased-noise conditions.~~

3.2 ~~FiLM conditioning~~

~~Feature-wise Linear Modulation (FiLM) [31] conditions each GNN layer on an external context. A small generator network maps the syndrome-weight estimate $\hat{p} = \|\mathbf{s}\|_1/(m_x + m_z)$ to per-layer affine parameters $(\gamma^{(\ell)}, \beta^{(\ell)}) = \mathrm{FiLMGen}(\hat{p})$, modulating hidden activations as $\tilde{\mathbf{h}}^{(\ell)} = \gamma^{(\ell)} \odot \mathbf{h}^{(\ell)} + \beta^{(\ell)}$. This conditioning pathway is designed to let a single model respond to global changes in the noise level without retraining. It should not be confused with spatial calibration: $\hat{p}$ is one scalar shared by the entire graph and therefore contains no direct per-qubit reliability information. Section 5.2 establishes a narrower result: uniform rescaling of the channel prior cannot change the hard decisions of the underlying unclipped min-sum BP decoder. That invariance does not, by itself, imply that a GNN cannot use a global scalar as contextual information.~~

3.2 ~~Neural BP and interleaved pipeline~~

~~Differentiable BP generalises (2)–(3) with learnable per-iteration scaling weights $w_{\mathrm{ch}}^{(t)}, w_{\mathrm{vtc}}^{(t)}, w_{\mathrm{ctv}}^{(t)} = \mathrm{softplus}(w_{\mathrm{raw}}^{(t)})$ (softplus ensures positivity; $w_{\mathrm{raw}} = 0$ at initialisation gives $w \approx \ln 2 \approx 0.693$, recovering near-standard BP).~~

~~Interleaved GNN-BP Decoding~~

~~$\mathrm{Decoding}(\gamma, \beta) \leftarrow \mathrm{FiLMGen}(\hat{p})$~~
~~$\Delta^{(1)} \leftarrow \textsc{TannerGNN}(\mathrm{A}, \mathbf{s}, \gamma,$~~
~~$\mathrm{A}' \leftarrow \mathrm{A} + \Delta^{(1)}$~~
~~$\mathrm{A}_{\mathrm{mid}} \leftarrow \mathrm{NeuralBP}(\mathrm{A}', \mathbf{H}, T_1)$~~
~~$p_i^{\mathrm{mid}} \leftarrow \sigma(-\mathrm{A}_{\mathrm{mid},i})\ \forall i$~~
~~$\Delta^{(2)} \leftarrow \textsc{TannerGNN}(\mathrm{A}_{\mathrm{mid}}, p^{\mathrm{mid}}, \mathbf{s}, \gamma, \beta)$~~
~~$\mathrm{A}'' \leftarrow \mathrm{A}_{\mathrm{mid}} + \Delta^{(2)}$~~
~~$\mathrm{A}_{\mathrm{final}} \leftarrow \mathrm{NeuralBP}(\mathrm{A}'', \mathbf{H}, T_2)$~~
~~$\hat{e}_i \leftarrow \mathbb{1}[\mathrm{A}_{\mathrm{final},i} < 0]\ \forall i$~~

~~Algorithm 1 gives the interleaved pipeline. The GNN intervenes twice: before BP to correct channel priors, and after T_1 iterations using the intermediate posterior information. Qubits with $p_i^{\mathrm{mid}} \approx 0.5$ have low posterior confidence and can receive a second learned correction. Both correction stages are trained jointly through the differentiable BP unroll.~~

~~Interleaved GNN-BP pipeline. The GNN applies two successive LLR corrections: one before BP initialisation and one after T_1 iterations, using the mid-decoding posterior as an additional input. Both stages are trained end-to-end through the BP unroll.~~

3.2 Training

The model is trained as a per-qubit classifier of the sampled Z error; BP is not part of its loss. For a shot with prior LLR ~~$\ell = \ln((1-p)/p)$~~ ℓ on every qubit, the GNN outputs a correction Δ_j for each qubit j, giving the corrected error probability $q_j = \sigma(-(\ell + \Delta_j))$~~. The model of Sec. 5.3 is trained with~~, and the loss is the focal loss [32]

Table 2: Architecture and training configuration of the TannerGNN model of Secs. 5.3 and 5.4.

GNN architecture	
Layers L / width d_h	3 / 32
Edge types	2 ($\mathbf{H}_x$, $\mathbf{H}_z$)
Node update	GRU, residual, layer norm
Node feature dim.	4
FiLM / attention	not used
Parameters	39,329
Training	
Data	8,000 shots, sinusoidal drift, $p \in [0.005, 0.045]$
Prior feature ℓ	$\ln((1-p)/p)$, true p of the slot, both components
Loss	focal, $(\alpha, \gamma) = (0.25, 2.0)$, on $q_j = \sigma(-(\ell + \Delta_j))$ vs. e_z; no syndrome term
Optimizer	AdamW, lr 2×10^{-3}, weight decay 10^{-4}
Batch size / epochs	64 / 20, gradient clipping 1.0
Checkpoint	fewest failures on 4,000 validation shots
Training seeds	5
BP ~~in training loop~~ for model selection	flooding min-sum, scaling 0.8, 10 iter.; not in the loss
Reference BP-OSD baseline	
Implementation	`ldpc` 2.4.1 `BpOsdDecoder`
`bp_method`	`ms` (min-sum)
`ms_scaling_factor`	0.8
`schedule`	`serial` (layered)
`max_iter`	100
OSD (where used)	`osd_cs`, order 10

against the true Z error e_z,

$$\mathcal{L} = -\frac{1}{Bn} \sum_{b=1}^{B} \sum_{j=1}^{n} \alpha_t \, (1 - q_t)^{\gamma} \log q_t, \qquad (7)$$

where $q_t = q_j$ and $\alpha_t = \alpha$ if $e_{z,j} = 1$, and $q_t = 1 - q_j$ and $\alpha_t = 1 - \alpha$ otherwise, with $\alpha = 0.25$ and $\gamma = 2$; no syndrome-consistency term is used. The computation graph is therefore $(\ell, \mathbf{s}) \to$ TannerGNN $\to \Delta \to q \to \mathcal{L}$, and gradients reach the GNN without passing through BP. Ten-iteration flooding BP enters training only through model selection: after each of the 20 epochs the corrected priors $\ell + \Delta$ are decoded with that BP on 4,~~and the corrected LLRs $\ell + \Delta_j$ are passed to both CSS components~~000 validation shots, and the epoch with the fewest failures is kept. The training data are 8,000 shots from two sinusoidally drifting streams, $p(t) = p_0 + 0.015 \sin(2\pi t/500)$ with $p_0 \in \{0.02, 0.03\}$, so the training rates span $p \in [0.005, 0.045]$. ~~The FiLM and interleaved variants of Secs. ?? ?? are trained on data regenerated each epoch with $p \sim$ Uniform$[0.01, 0.10]$,~~ Table 2 summarises the configuration.

The target is the sampled error e_z, so the loss is not degeneracy-aware: an estimate that differs from e_z by a stabilizer is logically correct but is penalized like any other error. The extreme case is an error that covers exactly half of a weight-six stabilizer g. Then e_z and ~~add a syndrome-consistency term, the binary cross-entropy between the measured syndrome and $\frac{1}{2}(1 - \prod_{j \in \mathcal{N}(c)}(1 - 2q_j))$ for each check c, with weight 0.1. Table 2 summarises the configuration.~~ $e_z \oplus g$ have the same weight and the same syndrome, so they are equally likely and the network receives the same input for both; the loss rewards treating the six qubits alike and gives no incentive to prefer either half. Section 5.4 shows that such configurations are a major cause of BP failures. Degeneracy-aware objectives, such as that of Liu and Poulin [24], avoid this problem; we did not use one.

Remark 1 (One correction for two CSS components)**.** *The GNN outputs one correction Δ_i per qubit and is trained on Z errors only. In the experiments of Sec. 5 the same Δ_i is added to the prior ℓ of both components, although the two error types differ: X errors are $\eta = 20$ times rarer, and the matched LLRs differ by a factor $\Lambda_x/\Lambda_z \approx 1.94$ at $p = 0.04$. Section 5.3 measures the consequence and shows that adding Δ_i to the Z component only is better.*

4 Experimental Setup

4.1 Evaluation regimes

Two distinct evaluation regimes are used throughout this paper, and the distinction is central to interpreting the results:

Regime A (~~training-aligned~~ development baseline): The GNN is evaluated ~~against~~ with flooding BP, the ~~base decoder used during training~~decoder used to select its checkpoint. This comparison isolates the gain relative to ~~the decoderrepresented in the training computation graph~~that decoder.

Regime B (stronger ~~classical baseline~~reference decoder): The same learned correction is evaluated with ~~the~~ serial-schedule min-sum BP-

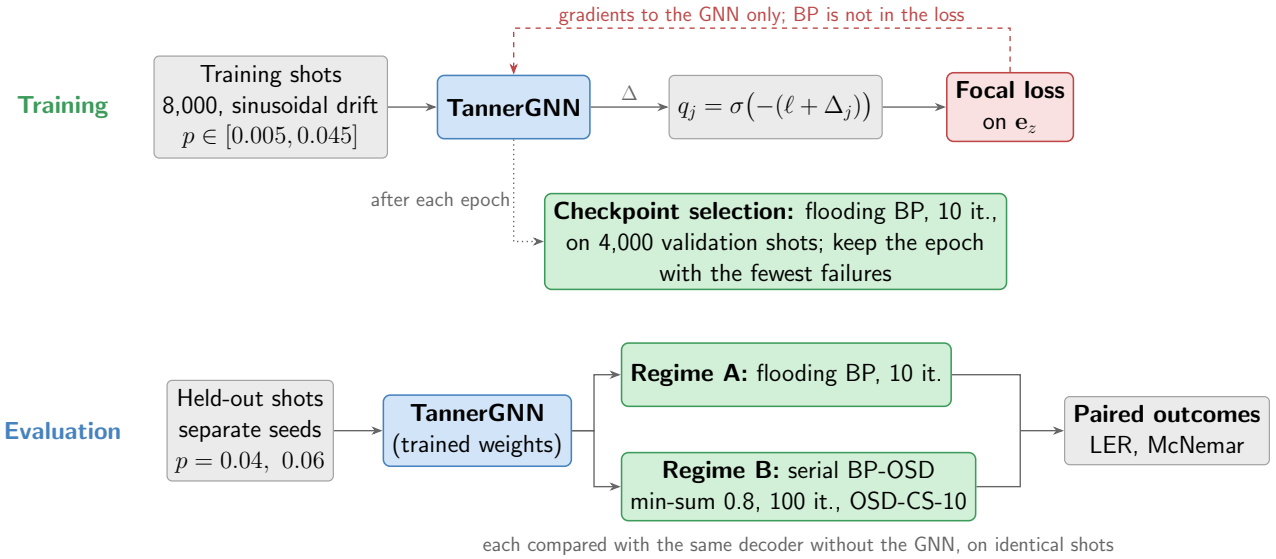


Figure 2: Training and evaluation pipeline of the GNN experiments. The GNN is trained ~~through~~ with a per-qubit focal loss, without BP in the loss; 10-iteration flooding BP ~~and~~ selects the checkpoint on validation shots. The model is evaluated on separately seeded held-out shots in Regime A (flooding BP) and Regime B (serial BP-OSD); each decoder receives identical syndrome batches to enable McNemar paired testing.

OSD ~~reference configuration~~(Table 2). This tests whether the correction also improves ~~the~~ a stronger classical decoder~~considered here;~~ ; BP-OSD serves here as a reference point, not as the best available decoder (Sec. 1.1). Plain BP-OSD uses a uniform prior in each component (p_z and p_x, or p in Table 6); the value of a uniform prior does not change its decisions, which we confirmed on all 20,000 test shots at $p = 0.04$. With the GNN, qubit i enters BP-OSD with prior error probability $\sigma(-(\ell + \Delta_i))$ in both components. Section 5.4 shows that ~~it does not~~ the correction does not help in the tested regime.

Unless otherwise stated, all experiments use Z-biased noise with $\eta = 20$.

4.2 Statistical protocol

Logical failure. For each shot let $\hat{e}_z$ and $\hat{e}_x$ be the decoder's hard decisions and $r_z = e_z \oplus \hat{e}_z$, $r_x = e_x \oplus \hat{e}_x$ the residual errors. A shot is a logical failure if (i) the estimate does not reproduce the measured syndrome, $\mathbf{H}_x \hat{e}_z \neq s_x$ or $\mathbf{H}_z \hat{e}_x \neq s_z$, or (ii) the syndrome is reproduced but the residual acts nontrivially on the code space, $\mathbf{L}_x r_z \neq \mathbf{0}$ or $\mathbf{L}_z r_x \neq \mathbf{0}$ over $\mathbb{F}_2$, where the rows of $\mathbf{L}_x$ ($\mathbf{L}_z$) span the $k = 12$ independent logical operators that detect Z (X) errors. Condition (i) is needed because a syndrome-inconsistent residual is not a logical operator, so its parity against fixed logical representatives is not meaningful; non-converged BP outputs therefore always count as failures. The X and Z components are combined by logical OR, so the reported LER is the probability that at least one of the 12 logical qubits suffers a logical error of any Pauli type in a shot. At circuit level the same rule is applied to the detector error model: a shot fails if the decoded fault set does not reproduce the detection events, or if it does and predicts the wrong observable flips.

For absolute LER comparisons we treat operating points with fewer than 100 observed logical-error events as preliminary; this is a reporting heuristic rather than a formal power calculation. Decoder pairs are evaluated on *identical* syndrome samples, enabling paired McNemar tests [33] (two-sided, with continuity correction). We report Wilson score 95% confidence intervals for LER estimates.

Comparison families. Each results table defines one family of paired tests: the ten GNN-BP vs. BP tests of Table 4 (two error rates, five seeds), the twenty GNN+OSD vs. BP-OSD tests of Table 6 (four error rates, five seeds), the nine oracle cells of Table 9, and the GNN comparisons of the drift experiment (Table 10). Within a family we control the familywise error rate with Holm's step-down procedure [34] at level 0.05 and say so where adjusted results

are reported; other p-values are unadjusted and exploratory. For emphasized comparisons we report the discordant counts n_{10} and n_{01}, and we never round a p-value to zero.

Training, validation and test data are drawn with distinct random seeds, so no evaluation shot shares a random stream with a training shot, and learned-decoder results are reported over five training seeds (initialisation and minibatch order). All training and evaluation used CPU only; no GPU was available for this work. One training epoch of the 39k-parameter model of Sec. 5.3 takes about 16 s.

5 Results

5.1 Decoder configuration sensitivity

Before presenting the GNN results, we quantify the sensitivity of the classical baseline to decoder configuration. This sensitivity sets the scale against which the learned-decoder results must be interpreted.

Table 3 illustrates how strongly decoder configuration changes the finite-length result. The original parallel product-sum BP-OSD configuration reduces LER from 0.0569 to 0.0398 relative to the corresponding plain BP, but remains far worse than the serial min-sum configurations. The reference serial min-sum BP-OSD records 34 logical errors in 10^5 shots, so its value is preliminary by the criterion of Sec. 4.2.

A separate $N = 100{,}000$-shot experiment (Fig. 3) changes one factor at a time. The flooding comparator is min-sum BP with scaling factor 0.8, at most 100 iterations, stopping at the first iteration whose hard decision satisfies the syndrome, with a parallel (flooding) schedule; the second decoder is identical except for the serial schedule; the third adds OSD-CS-10. The LER falls from 1.28×10^{-2} to 1.45×10^{-3} (8.8×, paired $n_{10}/n_{01} = 1138/3$) and then to 2.7×10^{-4} (5.4×, 118/0), about 47× overall.

These data do not identify a universally optimal configuration. They do show that schedule, update rule, iteration budget, and post-processing can dominate the measured finite-length LER. Learned-decoder comparisons must therefore specify and control these parameters.

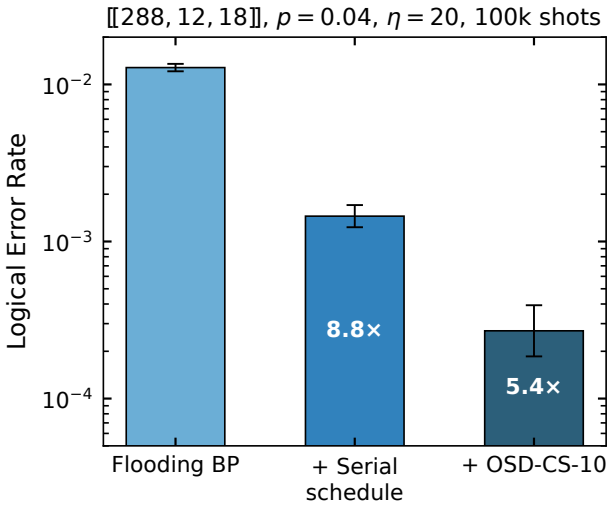


Figure 3: Baseline decomposition on $[\![288, 12, 18]\!]$, $p = 0.04$, $\eta = 20$, 100k identical shots. Min-sum BP (scaling 0.8, at most 100 iterations) with flooding and then serial schedule, then serial BP-OSD-CS-10. Wilson 95% CI error bars; labels give the factor gained at each step.

5.2 Uniform-prior scaling invariance of min-sum BP

Proposition 1 (Uniform positive-scaling invariance)**.** *Consider separate-CSS min-sum BP with no message clipping or offset. If all channel LLRs in one CSS component are multiplied by the same factor $a > 0$, then every BP message and posterior LLR is multiplied by a at every iteration. Consequently, all hard decisions are unchanged.*

Proof. At initialization, each variable-to-check message is proportional to the corresponding channel LLR and therefore scales by a. Assume all variable-to-check messages at iteration t have scaled by a. The syndrome factor $(-1)^{s_j}$ in Eq. (3) is unchanged, the sign of every message is unchanged because $a > 0$, and $\min_i a|x_i| = a \min_i |x_i|$. Hence every check-to-variable message also scales by a. Equation (2) then shows that the next variable-to-check messages scale by a. Induction establishes the claim for all iterations and for the final posterior LLRs. □

For a uniform channel prior, replacing a common positive LLR Λ by another positive value Λ' is equivalent to the scaling $a = \Lambda'/\Lambda$. The proposition explains why changing only the assumed uniform physical error rate cannot change hard decisions of ideal min-sum BP. Clipping, offsets, damping rules that depend on magnitude,

Table 3: Effect of BP-OSD configuration on $[\![288, 12, 18]\!]$, $p = 0.04$, $\eta = 20$, 100,000 identical shots, priors p_z and p_x of Eq. (4). 95% Wilson intervals in brackets.

Config	Schedule	Method	Iter	LER	Errors
Original BP	parallel	prod-sum	30	0.0569 [0.0555,0.0583]	5,689
Reference BP	serial	ms(0.8)	100	0.0017 [0.0015,0.0020]	174
Original BP-OSD	parallel	prod-sum	30	0.0398 [0.0386,0.0411]	3,983
Ref. BP-OSD	serial	ms(0.8)	100	0.00034 [0.00024,0.00048]	34

non-uniform priors, or sum-product updates can break the exact scaling argument.

We checked the proposition on 2,000 shots each. With unclipped min-sum BP (serial or flooding, 100 iterations), changing the assumed uniform rate from $p = 0.015$ to $p = 0.0286$ on $[\![288, 12, 18]\!]$, or from $p = 0.04$ to $p = 0.02$ on $[\![72, 12, 6]\!]$, leaves every hard decision unchanged (2,000/2,000). Two cases fall outside the proposition, and the invariance fails there: our PyTorch min-sum implementation clamps messages to ± 20 and changes 18 of 2,000 decisions on $[\![288, 12, 18]\!]$, and sum-product BP changes 25–43 of 2,000.

Remark 2 (What the proposition does and does not imply)**.** *The result applies to the BP pathway when adaptivity enters only by changing a uniform channel-prior magnitude. It does* not *establish that a global conditioning scalar is information-theoretically useless to a GNN: a learned model could combine a global noise estimate with local syndrome structure and change its nonlinear decision rule. The experiments below test a narrower hypothesis: spatially resolved per-qubit calibration can change the relative reliabilities seen by the decoder, whereas a single graph-wide scalar cannot provide that spatial information.*

5.3 GNN improvement over flooding BP

We evaluated the GNN ~~against~~ with flooding BP—the decoder ~~it was trained against~~ on used to select its checkpoint—on separately seeded held-out syndromes: 4,000 validation shots at $p = 0.04$ select the checkpoint over 20 epochs, and test sets of 20,000 shots at $p = 0.04$ (inside the training range) and 10,000 shots at $p = 0.06$ (outside it) are each scored once on that checkpoint. Training is repeated with five seeds.

Table 4 gives the primary learned-decoder result. Over five training seeds, GNN-BP reduces LER relative to flooding BP by $12.5\% \pm 0.6\%$ at $p = 0.04$ and by $10.6\% \pm 0.7\%$ at $p = 0.06$, outside the training range (McNemar $p < 10^{-36}$ for every seed, also after Holm correction over the ten tests). The gain comes mostly from convergence: shots on which BP does not reproduce the syndrome fall from 1,740 to 838–975 of 20,000. Flooding BP with 100 iterations and serial BP-OSD both remain better than GNN-BP. A gradient audit on the five checkpoints found non-zero signal through the network (input-to-readout gradient ratio ~~0.22~~0.63–2.19) and non-trivial LLR corrections (mean ~~$|\Delta_i| = 0.854$~~ $|\Delta_i| = 0.54$–1.01).

Table 4: GNN-BP vs. flooding BP on $[\![72, 12, 6]\!]$, $\eta = 20$, 20,000 ($p = 0.04$) and 10,000 ($p = 0.06$) held-out test syndromes, five training seeds (mean ± s.d.). n_{10} (n_{01}): shots only BP (only GNN-BP) fails. BP is flooding min-sum run for exactly 10 or 100 iterations; the last hard decision is used.

	$p = 0.04$	$p = 0.06$
BP, 10 iter.	~~0.1077~~ 0.1078	0.2867
GNN-BP, 10 iter.	0.0942 ± 0.0006	0.2564 ± 0.0019
relative reduction	$12.5\% \pm 0.6\%$	$10.6\% \pm 0.7\%$
n_{10} / n_{01}	313–370 / 64–95	343–394 / 66–80
BP, 100 iter.	~~0.0863~~ 0.0864	0.2403
Serial BP-OSD	0.0707	0.2138

Table 5: Failures by CSS component on $[\![72, 12, 6]\!]$, $p = 0.04$, $\eta = 20$, 20,000 held-out shots (counts). Z: Z errors decoded on $\mathbf{H}_x$; X: X errors decoded on $\mathbf{H}_z$; Either: the failure criterion of Sec. 4.2. GNN rows give the range over five training seeds.

Decoder	Z	X	Either
BP, 10 iter.	2,155	0	2,155
GNN-BP, Δ on both	1,845–1,899	7–54	1,874–1,906
GNN-BP, Δ on Z only	1,845–1,899	0	1,845–1,899
BP, 100 iter.	1,727	0	1,727
Serial BP-OSD	1,414	0	1,414
GNN + BP-OSD	1,439–1,451	0–1	1,439–1,451

The two CSS components. Table 5 separates the failures at $p = 0.04$ by component.

All failures of BP, 100-iteration BP and BP-OSD are in the Z component: at $p_x = 0.0019$ none of them fails on the X component in 20,000 shots. The whole GNN gain is therefore in the Z component. Adding Δ to the X component as well, as in Table 4, gains nothing there and creates 7–54 X-component failures. Adding it to the Z component only removes them and gives a reduction of $13.3\% \pm 0.9\%$ ($10.8\% \pm 0.7\%$ at $p = 0.06$). This variant was chosen after seeing the test results, so we keep the pre-specified setting in Table 4. Because BP makes no X-component errors at these operating points, a separately trained X head would have nothing to correct, and leaving the X prior unchanged is the appropriate choice. With BP-OSD the X component is essentially unaffected (at most one extra failure at $p = 0.04$), so the comparisons below do not depend on this choice.

A fixed-dataset capacity test provides an additional diagnostic: a GNN trained on 1,000 fixed syndromes for 50 epochs reduced the failure count on those same syndromes from 105 to 93 (11%; McNemar $p = 0.0095$, $n_{10} = 15$, $n_{01} = 3$). This experiment tests the ability of the architecture and optimizer to fit the supplied examples; it is not a generalization result. The ~~interleaved architecture of Sec. ??~~ is FiLM and interleaved variants (Appendix A) are not evaluated in this work.

5.4 ~~Training~~Development/evaluation baseline mismatch

Although the GNN improves flooding BP, the same correction does not improve the serial min-sum BP-OSD reference configuration. Note that Tables 4 and 6 report different evaluation regimes (Regime A and Regime B of Sec. 4.1) on different syndrome sets, and their LER values are not directly comparable~~.~~.

Table 6 and Fig. 4 show that the GNN-corrected decoder never improves on serial BP-OSD: ~~at $p = 0.02$~~ no seed is significantly better at any error rate. Without correction, one seed of five is significantly worse ~~, and from $p = 0.03$ on~~ at $p = 0.02$ and 14 of the 15 seed–error-rate pairs are significantly *worse* from $p = 0.03$ on; after Holm correction over the 20 tests, 11 of these 15 remain, including all five seeds at $p = 0.05$. The same ~~checkpoint improves~~ checkpoints improve flooding BP, so these data show that the learned correction does not transfer unchanged to the stronger decoder.

Table 6: ~~GNN-BP~~ GNN+OSD vs. serial BP-OSD on $[\![72, 12, 6]\!]$, $\eta = 20$, 20,000 identical shots per p, five training seeds (ranges over seeds). n_{10} (n_{01}): shots only BP-OSD (only GNN+OSD) fails. Last ~~column~~two columns: seeds for which GNN+OSD is significantly worse (McNemar $p < 0.05$), unadjusted and after Holm correction over the 20 tests. No seed is significantly better.

~~p Serial BP-OSD GNN+OSD Seeds worse 0.020 0.0093 0.0095–0.0102 1/5 0.030 0.0297 0.0306–0.0317 5/5 0.040 0.0676 0.0688–0.0709 4/5 0.050 0.1354 0.1384–0.1418 5/5~~

p	BP-OSD	GNN+OSD	n_{10} / n_{01}	Seeds worse unadj.	Holm
0.020	0.0093	0.0095–0.0102	3 / 7–20	1/5	1/5
0.030	0.0297	0.0306–0.0317	21–28 / 46–64	5/5	4/5
0.040	0.0676	0.0688–0.0709	52–62 / 84–120	4/5	2/5
0.050	0.1354	0.1384–0.1418	94–121 / 169–222	5/5	5/5

The GNN ~~was optimized through a~~ checkpoint was selected for 10-iteration ~~flooding-BP computation graph~~ flooding BP and is then inserted ahead of a decoder with a different schedule, iteration budget, and OSD post-processing. ~~The experiment therefore does not establish that LLR corrections learned through the flooding-BP computation graph remain useful after the schedule, iteration budget, and post-processing are changed.~~

A paired failure-set analysis (Table 7) shows that the ~~two decoders do not fail on disjoint sets : every shot that 100-iteration BP fails is also failed by 10-iteration BP ,~~ failure sets of BP and BP-OSD overlap strongly but are not nested: 1,319 shots fail under both, 836 only under BP and 95 only under BP-OSD, so 93% of BP-OSD's failures are BP failures and 61% of BP's failures are BP-OSD failures. The 100-iteration BP failures, by contrast, are a strict subset of the 10-iteration ones. What separates ~~them~~ BP from BP-OSD is mainly convergence: 81% of BP failures are shots on which BP does not reach a syndrome-consistent estimate, whereas BP-OSD always returns one. Of the 313–370 shots the GNN fixes for BP, 91–93% are non-converged BP shots and only 31–36% are shots that BP-OSD fails, so the GNN's gains lie mostly where BP-OSD already succeeds. When the same corrections are passed to BP-OSD they reshuffle its failures (237–282 newly failed vs. 206–254 rescued shots), almost entirely among shots on which BP itself fails. The same pattern holds at

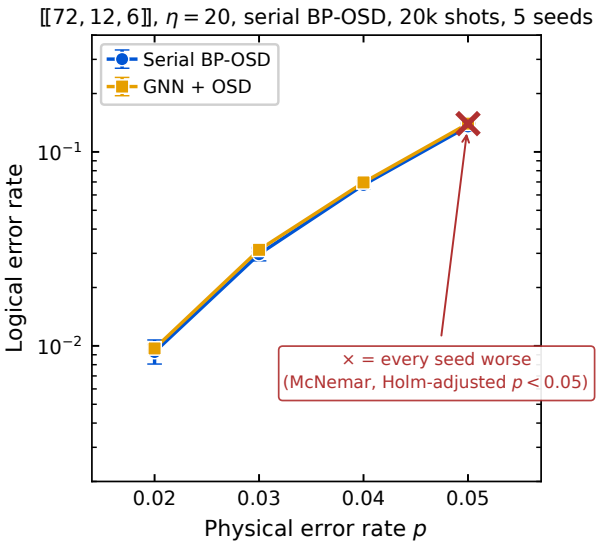


Figure 4: GNN+OSD vs. serial BP-OSD LER on $[\![72, 12, 6]\!]$, $\eta = 20$, 20,000 shots per point. BP-OSD: Wilson 95% intervals; GNN+OSD~~is the~~: mean over five training seeds, error bars give the range over seeds. ~~Red~~ The red $\times$ marks the operating ~~points~~ point where every seed is ~~statistically~~ significantly worse ~~(McNemar $p < 0.05$). Wilson 95% CI error bars~~after Holm correction.

$p = 0.06$, where 96% of BP-OSD failures are BP failures (and 71% the other way). Training directly against the target decoder, or against a differentiable surrogate of its post-processing stage, remains the appropriate test of whether learned augmentation can improve BP-OSD.

Which failures the GNN repairs. To see what kind of failure the correction repairs, we traced flooding BP for 100 iterations on the $p = 0.04$ test set (Table 8). Of the 2,155 BP failures, 415 reach a syndrome-consistent estimate in the wrong logical class. The other 1,740 do not reproduce the syndrome after 10 iterations: 78 had reproduced it earlier and moved away, 1,030 reproduce it later, 631 never reproduce it and keep changing their hard decision, and one is stuck. Stopping each shot at its first syndrome-consistent iteration would change the BP failure count only from 2,155 to 2,146, so the fixed iteration budget is not a handicap of the baseline. The shots the GNN repairs are mostly slow ones (249–265 of 313–370), followed by oscillating (37–78) and wrongly converged (22–32) shots: the correction mainly speeds up convergence.

A recurring structure among the failures is a symmetric stabilizer configuration [9]: the Z error covers exactly three qubits of a weight-six Z stabilizer g, so that e_z and $e_z \oplus g$ have the same weight and syndrome and are logically equivalent. Such configurations occur in 3.5% of the test shots but in 28% of BP failures; BP fails on 87% of the shots that contain one and on 8% of the others. Error weight does not explain this. On all 57 weight-three errors of this kind, BP, 100-iteration BP and GNN-BP with every seed fail, while BP-OSD fails on none (24 of 24 and none at $p = 0.06$); on the 4,466 weight-three errors without such a configuration BP fails on 2.4%. The GNN lowers BP's failure rate on symmetric configurations from 87% to 67–74%, all of it on configurations of weight four and more; on the weight-three ones it never helps, as expected from its loss (Sec. 3.2). OSD resolves the tie by construction, which is one reason why many BP failures are not BP-OSD failures.

Table 7: Paired failure sets on $[\![72, 12, 6]\!]$, $p = 0.04$, $\eta = 20$, 20,000 identical held-out shots. BP: 10-iteration flooding BP. GNN rows give the range over five training seeds. "Non-conv.": the estimate does not reproduce the syndrome.

Decoder	Failures	Non-conv.
BP, 10 iter.	2,155	1,740
BP, 100 iter.	1,727	700
Serial BP-OSD	1,414	0
GNN + BP, 10 iter.	1,874–1,906	838–975
GNN + BP-OSD	1,439–1,451	0
BP vs. BP-OSD failure sets		
Both fail	1,319	
Only BP fails / only BP-OSD fails	836 / 95	
BP fails, given BP-OSD fails	93%	
BP-OSD fails, given BP fails	61%	
BP-100 failures also BP failures	1,727 / 1,727	
What the GNN changes		
Shots GNN fixes for BP	313–370	
of which BP did not converge	91–93%	
of which BP-OSD also fails	31–36%	
BP-OSD failures also GNN+BP failures	87–88%	
GNN+BP-OSD vs. BP-OSD: newly failed / rescued	237–282 / 206–254	

5.5 Per-qubit calibration gap

Under spatially non-uniform per-qubit noise, each data qubit has rate $p_i = p\, e^{\sigma z_i - \sigma^2/2}$ with $z_i \sim \mathcal{N}(0, 1)$ drawn fresh every shot (capped at

Table 8: Classes of the BP failures on $[\![72, 12, 6]\!]$ at $p = 0.04$ (10-iteration flooding BP, traced for 100 iterations) and of the shots the GNN repairs (range over five seeds). "Met": the hard decision reproduces the syndrome. A symmetric configuration is a Z error covering three qubits of a weight-six Z stabilizer.

Class	BP failures	GNN repairs
Wrong class (syndrome met)	415	22–32
Transient (met before it. 10)	78	5–11
Slow (met by it. 100)	1,030	249–265
Oscillating (never met)	631	37–78
Stuck (never met)	1	0
Total	2,155	313–370
symmetric configuration	606	89–137

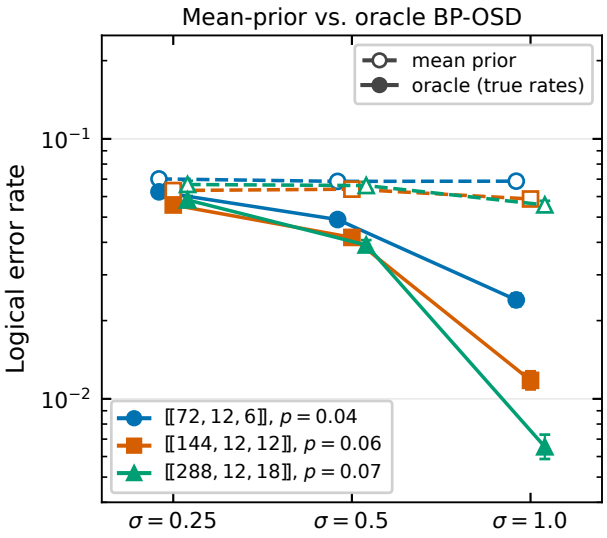


Figure 5: ~~Oracle calibration gap (percentage LER reduction from per-qubit vs. mean-prior~~ Mean-prior and oracle serial BP-OSD ~~)~~LER with Wilson 95% intervals, ~~50k~~ 50,000 identical shots per point. ~~Operating points:~~ Each code is shown at its own operating point ($p = 0.04, 0.06, 0.07$~~respectively~~); ~~cross-code comparisons are partially confounded by the differing~~ because code size and operating ~~points~~point change together, the figure does not show a distance trend.

0.5), so the mean rate is p for every σ; the Z and X rates of qubit i are $p_i\eta/(\eta+1)$ and $p_i/(\eta+1)$. A decoder knowing the true per-qubit rates (serial min-sum BP-OSD-CS-10 with the true p_i as priors) substantially outperforms a decoder using the uniform mean prior p.

Table 9 and Fig. 5 show a clear increase with noise contrast σ for each tested code. The reported percentage reduction is also larger for the larger-code rows, reaching 88% on $[\![288, 12, 18]\!]$ at $\sigma = 1.0$; however, the physical error rate is changed across codes ($p = 0.04$, 0.06, and 0.07). Consequently, these data do not isolate code distance as the cause of the cross-code trend. At a common operating point ($p = 0.04$, $\sigma = 1$) the reductions are 65%, 85% and 88% for the three codes, but the two larger codes record only 278 and 17 mean-prior failures there, so this controlled comparison is preliminary.

This experiment isolates the value of synthetic per-qubit rate side information; it does not require machine learning. We use the true-rate decoder as an *oracle reference*, not as a formal upper bound on all possible decoders. The comparison quantifies how much the tested BP-OSD implementation can benefit when the realized per-qubit rates are revealed. ~~The gap is *not* recoverable from a single syndrome shot~~under When the field is redrawn independently for every shot, these rates are not identifiable from that shot's syndrome: averaged over the field, the errors of each CSS component are independent across qubits with the uniform rates p_z and p_x, exactly as under uniform noise, so a component's syndrome carries no information about the realized field. (The field does correlate the X and Z errors of a qubit, which a joint decoder could exploit; ours decode the components separately.) A syndrome-only decoder could still outperform mean-prior BP-OSD by being a better decoder for i.i.d. ~~per-qubit noise: with a fresh field redrawn every shot, the uniform mean prior is Bayes-optimal within the class of estimators that observe only~~the syndrome, ~~so no syndrome-conditioned GNN can do better than the mean prior in this regime.~~ noise, but such a gain would not come from estimating the p_i.

5.6 Drift adaptation

~~Under persistent OU drift ($\sigma = 1.0$~~We test whether per-qubit rate estimates help BP-OSD when the rates drift. On $[\![72, 12, 6]\!]$ with mean rate $p = 0.04$ and $\eta = 20$, each data qubit i has a log-rate offset $x_i(t)$ that follows a discrete Ornstein–Uhlenbeck (AR(1)) process, independently across qubits,

$$x_i(t+1) = \rho\, x_i(t) + v\, \xi_i(t), \qquad \rho = \sqrt{1 - v^2/\sigma^2}, \tag{8}$$

with $\xi_i(t) \sim \mathcal{N}(0,1)$, $x_i(0) \sim \mathcal{N}(0, \sigma^2)$ and $\sigma = 1$, so the stationary spread of the log-rates is σ for

Table 9: Oracle gap: ~~percentage~~ LER ~~reduction from per-qubit over~~ of mean-prior and per-qubit (oracle) serial BP-OSD with Wilson 95% intervals, 50,000 identical shots per ~~cell, at $p = 0.04, 0.06, 0.07$ for~~ row; the three codes are at different operating points. ~~All McNemar $p < 10^{-20}$~~Reduction: $1 - \mathrm{LER}_{\mathrm{oracle}}/\mathrm{LER}_{\mathrm{mean}}$. ~~Parenthetical counts are n_{10} (n_{01}): shots only the~~ mean-prior ~~error events~~decoder (only the oracle) fails. All nine McNemar $p < 10^{-22}$, also after Holm correction.

Code (d)	~~$\sigma = 0.25$~~ p	~~$\sigma = 0.5$~~ σ	~~$\sigma = 1.0$~~ Mean prior	Oracle	Rec
$[\![72, 12, 6]\!]$ (6)	~~+11% (3508)~~ 0.04	~~+29% (3432)~~ 0.25	~~+65% (3438)~~ 0.0702 [0.0680,0.0724]	0.0626 [0.0605,0.0647]	
		0.5	0.0686 [0.0665,0.0709]	0.0489 [0.0471,0.0509]	
		1.0	0.0688 [0.0666,0.0710]	0.0240 [0.0227,0.0254]	
$[\![144, 12, 12]\!]$ (12)	~~+12% (3165)~~ 0.06	~~+35% (3203)~~ 0.25	~~+80% (2936)~~ 0.0633 [0.0612,0.0655]	0.0557 [0.0537,0.0577]	
		0.5	0.0641 [0.0619,0.0662]	0.0418 [0.0401,0.0436]	
		1.0	0.0587 [0.0567,0.0608]	0.0118 [0.0109,0.0128]	
$[\![288, 12, 18]\!]$ (18)	~~+13% (3338)~~ 0.07	~~+41% (3307)~~ 0.25	~~+88% (2786)~~ 0.0668 [0.0646,0.0690]	0.0582 [0.0562,0.0603]	
		0.5	0.0661 [0.0640,0.0684]	0.0391 [0.0374,0.0408]	
		1.0	0.0557 [0.0537,0.0578]	0.0065 [0.0059,0.0073]	

every volatility v. The volatility is the per-shot standard deviation of the log-rate innovation: $v = 0$ is a static field, and $v = 0.1$, 0.35 and 0.5 give correlation times $-1/\ln\rho$ of about 200, 15 and 7 shots. The rate of qubit i at shot t is $p_i(t) = \min(0.5,\, p\,e^{x_i(t)-\sigma^2/2})$, whose stationary mean is p, and each shot draws independent Z and X flips with probabilities $p_i\eta/(\eta+1)$ and $p_i/(\eta+1)$.

All decoders are serial min-sum BP-OSD (Table 2) with per-qubit priors $\hat{p}_i\eta/(\eta+1)$ and $\hat{p}_i/(\eta+1)$; they differ only in the rate estimate $\hat{p}_i(t)$. The *oracle* uses the true $p_i(t)$. *Stale* uses $p_i(0)$, a perfect calibration taken at the start of the stream and never updated. *Mean* uses the uniform rate p. *Empir* counts the Z flips n_i in its own decoded estimates over the previous $W = 32$ ~~-round window, $[\![72, 12, 6]\!]$, $p = 0.04$, 6~~shots and shrinks the frequency toward p with $K = 16$ pseudo-shots, $\hat{p}_i = \frac{\eta+1}{\eta}\,(n_i + Kp\eta/(\eta+1))/(W+K)$. *GNN* is a TannerGNN rate estimator with the architecture of Table 2 that maps the violated-check frequencies of the previous W syndromes to an estimate $\hat{x}_i(t)$ and sets $\hat{p}_i = \min(0.5,\, p\,e^{\hat{x}_i-\sigma^2/2})$; it sees syndromes only, never decoder outputs or true rates. It is trained by mean-squared error on 20,000 ~~shots), the GNN beats stale calibration only at the highest tested drift volatility~~windows drawn at the four volatilities, with five training seeds; for each seed the epoch with the lowest validation error out of 30 is kept. Every decoder sees the current syndrome. Each volatility is evaluated on 20 independent streams of 2,532 shots, whose first 32 shots are warm-up and not scored, giving 50,000 scored shots. Shots in one stream share a rate field, so pooled McNemar tests overstate the evidence; we therefore compare decoders by a sign test over the 20 streams, with Holm correction over the eight GNN comparisons.

Table 10 and Fig. 6 show ~~that the GNN beats stale calibration at vol$= 0.5$ ($p = 0.005$, 158~~ three things. First, a per-qubit calibration is valuable only while it is current. Without drift the stale calibration is the oracle (LER 0.0210, against 0.0645 for the mean prior), but at $v = 0.1$, where the rates decorrelate over about 200 shots, it is already worse than the mean prior (0.0747 vs. ~~111 discordant pairs), but is beaten by~~ 0.0655). Second, tracking the rates recovers much of the gap: the non-learned ~~empirical frequency estimator at every volatility level (GNN~~ frequency estimator reaches 0.0333 without drift and 0.0358 at $v = 0.1$, and its advantage shrinks as the correlation time approaches the window length (0.0601 vs. ~~empir: $p = 2.7 \times 10^{-8}$ at vol$= 0.5$, GNN worse)~~0.0678 for the mean prior at $v = 0.5$). Third, the GNN estimator beats stale calibration at every nonzero volatility and the mean prior at every volatility (on all 20 streams up to $v = 0.35$ and on 17 of 20 at $v = 0.5$), but it is worse than the frequency estimator on all 20 streams at every volatility; all eight GNN comparisons remain significant after Holm correction. The five training seeds give nearly identical results. Under independent per-qubit drift, ~~direct per-qubit syndrome frequency counting captures most~~ counting each qubit's decoded errors therefore captures more of the

Table 10: Drift experiment on $[\![72, 12, 6]\!]$, $p = 0.04$, $\eta = 20$, $\sigma = 1$, $W = 32$: LER ~~under~~ of serial BP-OSD with different per-qubit ~~OU drift by volatility~~ rate estimates, 50,000 scored shots (~~vol~~20 streams) per volatility v. ~~"Empir"~~GNN: ~~non-learned per-qubit frequency~~ range over five training seeds. Last two columns: streams on which the seed-averaged GNN has fewer / more failures than stale calibration and than the empirical estimator.

~~vol~~ v	Oracle	Stale	Mean	Empir	GNN
0.00	~~0.033~~ 0.0210	~~0.033~~ 0.0210	~~0.057~~ 0.0645	~~0.064~~ 0.0333	~~GNN worse, $p = 5\times10^{-36}$~~ 0.0428–0.0433
0.10	~~0.031~~ 0.0223	~~0.038~~ 0.0747	~~0.057~~ 0.0655	~~0.067~~ 0.0358	~~GNN worse, $p = 8\times10^{-30}$~~ 0.0474–0.0477
0.35	~~0.005~~ 0.0232	~~0.027~~ 0.0821	~~0.038~~ 0.0682	~~0.050~~ 0.0541	~~GNN worse, $p = 5\times10^{-19}$~~ 0.0602–0.0607
0.50	~~0.004~~ 0.0243	~~0.052~~ 0.0819	~~0.035~~ 0.0678	**~~0.044~~** 0.0601	~~GNN beats stale, $p = 0.005$~~ 0.0648–0.0651

available signal ~~, leaving little room for a learned estimator. Whether spatial correlation structure between qubits would shift this result~~ than the learned syndrome-based estimator; whether spatially correlated drift would change this is an open question.

5.7 Circuit-level evaluation

Circuit and noise model. We simulate a Z-basis memory experiment on $[\![72, 12, 6]\!]$ with 6 rounds of syndrome extraction (72 data qubits, 36 X-check and 36 Z-check ancillas). Data qubits are prepared in $|0\rangle$; each round measures all X-checks and then all Z-checks with dedicated ancillas, using CNOT layers from a greedy conflict-free schedule; the data qubits are finally measured in the Z basis. Detectors are (i) the 36 Z-checks of round 1, which are deterministic for the prepared state, (ii) the XOR of consecutive measurements of every check in rounds 2–6, and (iii) the XOR of the last Z-check outcome with the parity of the measured data qubits in its support, giving 432 detectors and 12 logical observables. ~~Noise is a biased single-qubit Pauli channel with $p_X = p/(\eta+1)$, $p_Z = p\eta/(\eta+1)$, $p_Y = 0$, $\eta = 20$, applied to every qubit after each~~ The noise, summarized in Table 11, is a custom biased Pauli model, not standard circuit-level depolarizing noise: every one- and two-qubit gate ~~and to every idle qubit at each time step; measurements and resets flip with probability p~~ is followed by an independent single-qubit channel on each qubit it acts on, with mutually exclusive X and Z faults and no Y fault, so a two-qubit gate has no correlated two-qubit faults. Under spatially non-uniform noise each qubit q (data and ancilla) has rate $p_q = p\, e^{\sigma z_q - \sigma^2/2}$, $z_q \sim \mathcal{N}(0,1)$, $\sigma = 1$, on all of its fault locations, so the mean rate is p as in Sec. 5.5.

Table 11: Circuit noise for a qubit q with rate p_q ($p_q = p$ under uniform noise), $\eta = 20$. Each line acts independently on every qubit concerned.

Location	Fault
After reset to $\|0\rangle$	X, prob. p_q
After H or CNOT (each qubit)	X, prob. $p_q/(\eta+1)$, or Z, prob. $p_q\eta/(\eta+1)$
Idle qubit, each time step	as after a gate
Before Z measurement	outcome flip, prob. p_q

Decoders. ~~We decode the~~ Stim [35] converts each noisy circuit into a detector error model (DEM) with 4,068 fault mechanisms ~~)~~ on 432 detectors. Each mechanism collects the circuit faults that flip the same detectors and observables, and its prior is their combined probability; the exclusive X and Z outcomes of a channel are treated as independent faults (Stim's `approximate_disjoint_errors`). The mean-prior decoder uses the DEM of the circuit with every qubit at rate p, and the oracle uses the DEM of the circuit with each qubit at its true rate p_q; the two DEMs have the same mechanisms and differ only in their priors. Both are decoded with serial min-sum BP-OSD (scaling 0.8, 100 iterations, OSD-CS order 10), as in Table 2. ~~The mean-prior decoder uses the uniform rate p; the oracle uses the true p_q.~~ We draw 20 rate profiles at $p \le 0.003$ and 40 at $p = 0.01$, with 500 shots per profile. Failures follow Sec. 4.2.

Knowing the per-qubit rates also helps at circuit level (Table 12, Fig. 7). At $p = 0.01$ the oracle reduces the logical error rate from 1.35% to 0.60% (55%; McNemar $p = 4.8 \times 10^{-25}$), and it is better on 38 of 40 independent rate profiles (sign test $p = 1.5 \times 10^{-9}$), so the result is not driven by a few extreme profiles. At $p = 0.001$ and $p = 0.003$ neither decoder fails in 10^4 shots,

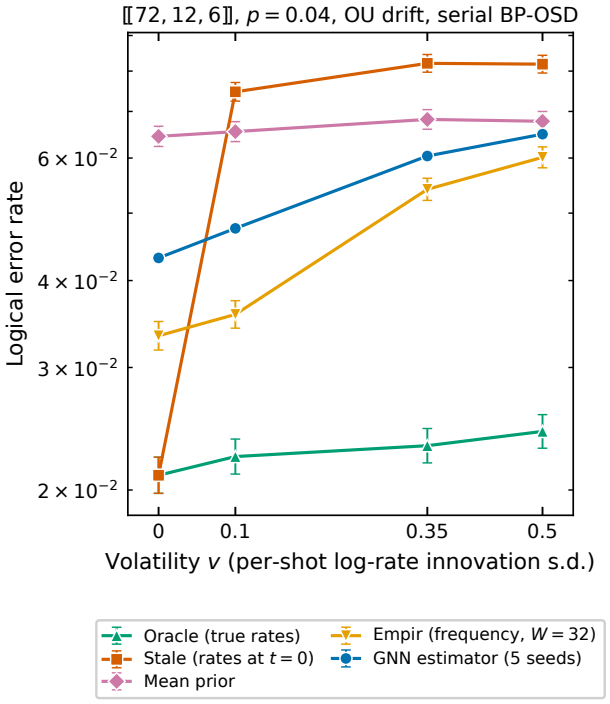


Figure 6: LER ~~vs.~~ of serial BP-OSD under per-qubit OU drift against the volatility ~~for four decoder configurations~~ v on $[\![72, 12, 6]\!]$, $p = 0.04$, ~~$\sigma_{\mathrm{OU}} = 1.0, 6$~~$\sigma = 1$, ~~000 shots~~for five rate estimates. ~~Star marks the single operating point (vol= 0.5) where GNN beats stale calibration ($p = 0.005$)~~Wilson 95% intervals; for the GNN~~is beaten by~~, the ~~empirical frequency estimator at every volatility level~~mean over five training seeds with the range over seeds.

so these operating points cannot resolve a gap. The circuit-level gap is somewhat smaller than the code-capacity gap on the same code (65% at $\sigma = 1$, Table 9), at a different operating point, but it does not vanish.

Table 12: Circuit-level oracle gap on $[\![72, 12, 6]\!]$ (6 rounds, $\eta = 20$, $\sigma = 1$). LER in percent with Wilson 95% intervals. n_{10} (n_{01}): shots only the mean-prior decoder (only the oracle) fails. ~~Zero failures in 10^4 shots means LER $< 3.8 \times 10^{-4}$ (95% upper bound).~~

~~p Shots Mean prior Oracle n_{10}/n_{01} p_{McN} 0.001 10,000 0 0 0/0 — 0.003 10,000 0 0 0/0 — 0.01 20,000 1.35% 0.60% 177/28 5×10^{-25}~~

p	Shots	Mean prior	Oracle	n_{10}	n_{01}	p_{McN}
0.001	10,000	0 [0,0.04]	0 [0,0.04]	0	0	—
0.003	10,000	0 [0,0.04]	0 [0,0.04]	0	0	—
0.01	20,000	1.35 [1.19,1.51]	0.60 [0.50,0.72]	177	28	5×10^{-25}

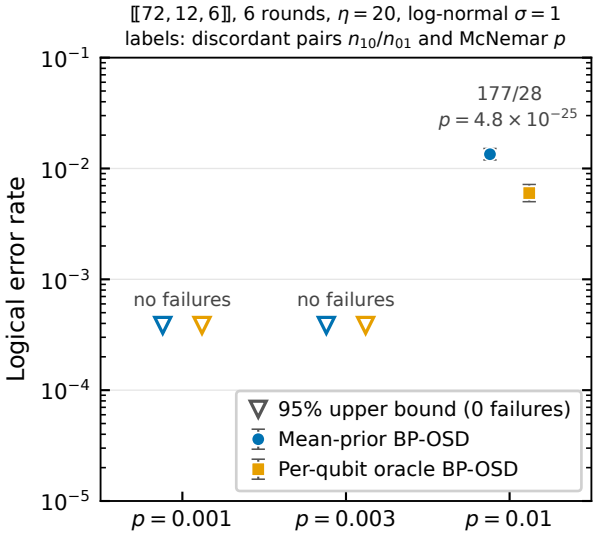


Figure 7: Circuit-level oracle gap on $[\![72, 12, 6]\!]$, 6 rounds, $\eta = 20$, $\sigma = 1$. Points: LER with 95% Wilson intervals; open triangles: 95% upper bound when no failures were observed. Labels: discordant pairs and McNemar p.

6 Discussion

6.1 The GNN's role and its limits

TannerGNN reduces the LER of the flooding BP decoder used ~~in its training loop~~ to select its checkpoint by $12.5\% \pm 0.6\%$ over five training seeds, and gradient diagnostics show that the network produces non-trivial corrections. The gain comes mainly from helping BP converge on shots where it converges slowly, and 100-iteration BP and serial BP-OSD remain stronger than the corrected decoder.

The Regime-B result is also informative: the learned correction does not transfer automatically to serial BP-OSD. This is consistent with the ~~training/evaluation mismatch : the failure sets of BP and~~ mismatch between the development and evaluation decoders: most BP-OSD ~~are nested~~failures are also BP failures, and the GNN's repairs fall mostly on shots BP-OSD already decodes (Table 7). The symmetric stabilizer configurations of Sec. 5.4 show the reverse case: failures that BP-OSD avoids by construction and that a bitwise loss cannot teach the GNN to repair. Training against BP-OSD directly is complicated by the non-differentiable OSD stage; differentiable surrogates, distillation, ~~or~~ training on the residual failures of the target decoder, or a degeneracy-aware loss [24] are possible alternatives.

6.2 Benchmark design implications

The baseline experiments identify two reporting requirements for learned QLDPC decoders. First, BP-OSD schedule, update rule, iteration count, and OSD configuration must be stated explicitly because they can change LER by more than two orders of magnitude on the same code. Table 2 records these settings for the present study.

Second, the ~~training and evaluation decoders~~ decoder used for model selection and the evaluation decoder must be distinguished. An LLR correction ~~learned through~~ selected for flooding BP need not remain beneficial when the schedule, iteration count, and OSD stage are changed. Comparisons across such decoder configurations should identify the mismatch rather than treating all BP-based baselines as equivalent.

6.3 Comparison with prior work

Astra [6] uses a learned Tanner-graph message-passing decoder and reports BB-code performance competitive with, and in several settings better than, BP-OSD under code-capacity depolarizing noise. This places our 12.5% flooding-BP gain in context: improvement over plain BP is not, by itself, the main novelty. Our contribution is instead the controlled comparison showing how that gain changes when the baseline and side information are strengthened.

Gong et al. [5] interleave differentiable BP stages with GNN corrections on the same sparse graph and show error-floor improvements over several post-processing approaches. Our interleaved variant (Appendix A) follows the same broad design principle, but its decoding performance is not evaluated here, so we do not claim a competitive result for that architecture.

Recent circuit-level work on BB codes also reports neural decoders that can outperform BP-OSD on the $[\![72, 12, 6]\!]$ code while retaining a runtime advantage, although scaling to $[\![144, 12, 12]\!]$ remains more challenging [25]. This comparison reinforces the need to evaluate learned decoders against a fully specified ~~competitive~~ strong non-learned baseline, not only against flooding BP.

Stein et al. [26] condition a neural QEC decoder on spatially resolved hardware calibration data using FiLM and evaluate on IBM repetition-code experiments. Their conditioning signal is per-qubit and calibration-derived, unlike the single ~~scalar $\hat{p}$ used in our initial model~~global rate p fed to our FiLM variant (Appendix A). This distinction is consistent with the oracle experiments of Sec. 5.5, which motivate spatial rather than purely global conditioning.

6.4 Towards spatially-resolved conditioning

The oracle-gap experiment (Sec. 5.5) shows that per-qubit rate side information can materially improve the tested BP-OSD decoder. Because operating point and code size co-vary, the current data do not establish how this value scales with distance. Proposition 1 also shows that a uniform rescaling of the min-sum BP prior cannot reproduce the effect of spatially varying reliabilities. These results motivate testing spatially resolved FiLM (SRCF), in which per-qubit parameters (γ_i, β_i) are conditioned on persistent hardware calibration data rather than on a single graph-wide scalar. Soft syndrome information [36] is another form of side information that iterative decoders can use. Such a decoder should be compared with the non-learned frequency estimator of Sec. 5.6 and evaluated in a circuit-level setting where the calibration signal remains identifiable. Neither condition has been established in the present experiments.

6.5 Limitations

(i) ~~All GNN variants were trained against~~ The evaluated GNN was trained with a per-qubit loss and selected for 10-iteration flooding BP~~. No variant was trained against~~; no model was trained or selected for serial BP-OSD, ~~which is the competitive classical baseline~~our reference decoder. Whether GNN augmentation of serial BP-OSD yields gains is untested.

(ii) GNN checkpoints exist only for $[\![72, 12, 6]\!]$; the larger codes were not trained due to CPU-only compute constraints. The null result in Regime B should not be extrapolated to the larger codes.

(iii) The interleaved and FiLM architectures (Appendix A) are described but their decoding performance is not evaluated.

(iv) ~~All training and evaluation used CPU only, which limited hyperparameter search and training duration.~~

(v) ~~The single per-qubit correction Δ_i is applied to both $A_{z,i}$ and$A_{x,i}$, which is suboptimal under $\eta = 20$ biased noise; separate~~ The correction is trained on Z ~~/~~errors and, in the headline setting, also added to the X ~~readout heads are the recommended fix~~component, where it only adds failures (Table 5); the Z-only setting was chosen after seeing the test results.

(vi) The training loss is not degeneracy-aware (Sec. 3.2).

(vii) Code-capacity noise omits measurement errors, leakage, and correlated faults present in real circuits, and the circuit model uses independent single-qubit channels rather than correlated two-qubit faults.

(viii) The drift experiment is synthetic: the rates of different qubits drift independently, with correlation times of 7 to 200 shots.

7 Conclusion

On $[\![72, 12, 6]\!]$, a Tanner-graph GNN ~~trained through~~ whose checkpoint is selected with 10-iteration flooding BP reduces the held-out LER of that decoder by $12.5\% \pm 0.6\%$ at $p = 0.04$ ~~and $\eta = 20$~~ over five training seeds, all of it in the Z component. The same ~~learned correction does not improve the~~ correction never improves serial min-sum BP-OSD ~~reference decoder and is significantly worse from $p = 0.03$ on; the shots it repairs are largely ones~~ and is often significantly worse; it mostly repairs shots on which BP converges slowly and which BP-OSD already decodes. Weight-three errors that cover half of a stabilizer defeat BP with or without the GNN but never BP-OSD, and the bitwise training loss cannot teach the network to break that tie. The demonstrated gain is therefore specific to the ~~training-aligned~~ flooding-BP baseline~~; the present experiments do not establish an improvement over competitive BB decoding more generally~~.

~~The baseline experiments show that~~ BP schedule, update rule, iteration budget ~~,~~ and OSD configuration ~~can change~~ change the finite-length LER by orders of magnitude~~. These parameters must be treated as part of~~ , so they belong to the decoder specification~~when learned and non-learned methods are compared. We also proved that uniform positive~~ . A uniform rescaling of the channel LLRs ~~leaves the hard decisions of ideal~~ cannot change min-sum ~~BP unchanged. A global prior magnitude alone therefore cannot provide spatial adaptation to heterogeneous qubit reliabilities, although a nonlinear learned model may still use a global noise estimate as context.~~

~~The synthetic oracle experiments identify per-qubit calibration as the more informative adaptive signal in the tested setting: revealing~~ decisions, whereas the true per-qubit rates ~~substantially improves~~ reduce BP-OSD~~under heterogeneous noise, and~~ 's LER by up to 88% under synthetic heterogeneous noise and by 55% under circuit-level noise~~it reduces BP-OSD's LER by 55% at $p = 0.01$. The current cross-code experiments do not isolate how this gain scales with distance. A direct~~ . Under per-qubit drift a stale calibration quickly loses its value, and a simple frequency estimator tracks the rates better than a learned one. A natural next test is ~~to condition a learned decoder~~ a decoder conditioned on persistent spatial calibration data~~and compare it against a fully specified serial~~ , trained with a degeneracy-aware loss, and compared with fully specified BP-OSD ~~baseline and the non-learned frequency estimator, using multiple training seeds and a noise model that preserves temporal calibration structure~~and the frequency estimator.

Acknowledgments

The authors thank the STAR Lab program at the University of Arizona for facilitating this research and providing computational resources.

Data and code availability

~~Source code, trained checkpoints, datasets and the analysis scripts~~ All code, data, trained checkpoints and per-shot outcomes behind every table and figure ~~except the drift experiment (Sec. 5.6)~~ are available at `https://github.com/kondshk/adaptive_gnn` (tag `paper-v5`). The script `make_tables.py` rebuilds every results table from the stored outcomes, and `requirements-lock.txt` lists the exact package versions. Code-capacity data are sampled directly with NumPy; circuit-level data are generated with Stim [35].

References

[1] Sergey Bravyi, Andrew W. Cross, Jay M. Gambetta, Dmitri Maslov, Patrick Rall, and Theodore J. Yoder. "High-threshold and low-overhead fault-tolerant quantum memory". Nature **627**, 778–782 (2024).

[2] IBM Quantum. "IBM lays out clear path to fault-tolerant quantum computing". IBM Quantum Computing Blog (2025).

[3] Theodore J. Yoder, Eddie Schoute, Patrick Rall, Emily Pritchett, Jay M. Gambetta, Andrew W. Cross, Malcolm Carroll, and Michael E. Beverland. "Tour de gross: A modular quantum computer based on bivariate bicycle codes" (2025). arXiv:2506.03094.

[4] Joschka Roffe, David R. White, Simon Burton, and Earl Campbell. "Decoding across the quantum LDPC code landscape". Physical Review Research **2**, 043423 (2020).

[5] Anqi Gong, Sebastian Cammerer, and Joseph M. Renes. "Graph neural networks for enhanced decoding of quantum LDPC codes". In 2024 IEEE International Symposium on Information Theory (ISIT). Pages 2700–2705. (2024).

[6] Arshpreet Singh Maan and Alexandru Paler. "Machine learning message-passing for the scalable decoding of QLDPC codes". npj Quantum Information **11**, 78 (2025).

[7] David J. C. MacKay, Graeme Mitchison, and Paul L. McFadden. "Sparse-graph codes for quantum error correction". IEEE Transactions on Information Theory **50**, 2315–2330 (2004).

[8] David Poulin and Yeojin Chung. "On the iterative decoding of sparse quantum codes". Quantum Information and Computation **8**, 987–1000 (2008).

[9] Nithin Raveendran and Bane Vasić. "Trapping sets of quantum LDPC codes". Quantum **5**, 562 (2021).

[10] Kao-Yueh Kuo and Ching-Yi Lai. "Exploiting degeneracy in belief propagation decoding of quantum codes". npj Quantum Information **8**, 111 (2022).

[11] Marc P. C. Fossorier and Shu Lin. "Soft-decision decoding of linear block codes based on ordered statistics". IEEE Transactions on Information Theory **41**, 1379–1396 (1995).

[12] Pavel Panteleev and Gleb Kalachev. "Degenerate quantum LDPC codes with good finite length performance". Quantum **5**, 585 (2021).

[13] Julien du Crest, Mehdi Mhalla, and Valentin Savin. "Stabilizer inactivation for message-passing decoding of quantum LDPC codes". In 2022 IEEE Information Theory Workshop (ITW). Pages 488–493. (2022).

[14] Julien du Crest, Francisco Garcia-Herrero, Mehdi Mhalla, Valentin Savin, and Javier Valls. "Check-agnosia based post-processor for message-passing decoding of quantum LDPC codes". Quantum **8**, 1334 (2024).

[15] Timo Hillmann, Lucas Berent, Armanda O. Quintavalle, Jens Eisert, Robert Wille, and Joschka Roffe. "Localized statistics decoding for quantum low-density parity-check codes". Nature Communications **16**, 8214 (2025).

[16] Sana Javed, Sergio Colet, Francisco Garcia-Herrero, Óscar Ruano, Juan Antonio Maestro, Bane Vasić, and Mark F. Flanagan. "Low-complexity syndrome-based linear programming decoding of quantum LDPC codes". IEEE Transactions on Quantum Engineering **7**, 2100319 (2026).

[17] Dimitris Chytas, Nithin Raveendran, and Bane Vasić. "Collective bit flipping-based decoding of quantum LDPC codes". IEEE Transactions on Communications **73**, 5566–5579 (2025).

[18] Asit Kumar Pradhan, Nithin Raveendran, Narayanan Rengaswamy, and Bane Vasić. "Linear time iterative decoders for hypergraph-product and lifted-product codes". IEEE Transactions on Information Theory **72**, 4618–4647 (2026).

[19] Tristan Müller, Thomas Alexander, Michael E. Beverland, Markus Bühler, Blake R. Johnson, Thilo Maurer, and Drew Vandeth. "Improved belief propagation is sufficient for real-time decoding of quantum memory" (2025). arXiv:2506.01779.

[20] Arshpreet Singh Maan, Francisco Miguel Garcia Herrero, Alexandru Paler, and Valentin Savin. "Decoding correlated errors in quantum LDPC codes". Nature Communications **17**, 3965 (2026).

[21] Francisco Garcia-Herrero, Javier Valls, Llanos Vergara-Picazo, and Vicente Torres. "Diversity methods for improving convergence and accuracy of quantum

error correction decoders through hardware emulation". Quantum **10**, 2071 (2026).

[22] Stasiu Wolanski and Ben Barber. "Ambiguity clustering: an accurate and efficient decoder for qLDPC codes" (2024). arXiv:2406.14527.

[23] Daniel Báscones, Francisco Garcia-Herrero, and Javier Valls. "Exploring the FPGA and ASIC design space of belief propagation and ordered statistics decoders for quantum error correction codes". EPJ Quantum Technology **12**, 140 (2025).

[24] Ye-Hua Liu and David Poulin. "Neural belief-propagation decoders for quantum error-correcting codes". Physical Review Letters **122**, 200501 (2019).

[25] John Blue, Harshil Avlani, Zhiyang He, Liu Ziyin, and Isaac L. Chuang. "Machine learning decoding of circuit-level noise for bivariate bicycle codes". Quantum **10**, 2149 (2026).

[26] Samuel Stein, Shuwen Kan, Chenxu Liu, Adrian Harkness, Sean Garner, Zefan Du, Yufei Ding, Ying Mao, and Ang Li. "Calibration-conditioned featurewise linear modulation decoders for low-latency decoding of quantum error correction evaluated on IBM repetition-code experiments". Physical Review Research **8**, 033194 (2026).

[27] A. Robert Calderbank and Peter W. Shor. "Good quantum error-correcting codes exist". Physical Review A **54**, 1098 (1996).

[28] Andrew M. Steane. "Error correcting codes in quantum theory". Physical Review Letters **77**, 793 (1996).

[29] P. V. Klimov et al. "Fluctuations of energy-relaxation times in superconducting qubits". Physical Review Letters **121**, 090502 (2018).

[30] Malcolm Carroll, Sami Rosenblatt, Petar Jurcevic, Isaac Lauer, and Abhinav Kandala. "Dynamics of superconducting qubit relaxation times". npj Quantum Information **8**, 132 (2022).

[31] Ethan Perez, Florian Strub, Harm de Vries, Vincent Dumoulin, and Aaron Courville. "FiLM: Visual reasoning with a general conditioning layer". In Proc. AAAI. Volume 32. (2018). arXiv:1709.07871.

[32] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. "Focal loss for dense object detection". In Proc. IEEE ICCV. Pages 2999–3007. (2017).

[33] Quinn McNemar. "Note on the sampling error of the difference between correlated proportions or percentages". Psychometrika **12**, 153–157 (1947).

[34] Sture Holm. "A simple sequentially rejective multiple test procedure". Scandinavian Journal of Statistics **6**, 65–70 (1979).

[35] Craig Gidney. "Stim: A fast stabilizer circuit simulator". Quantum **5**, 497 (2021).

[36] Nithin Raveendran, Javier Valls, Asit Kumar Pradhan, Narayanan Rengaswamy, Francisco Garcia-Herrero, and Bane Vasić. "Soft syndrome iterative decoding of quantum LDPC codes and hardware architectures". EPJ Quantum Technology **10**, 45 (2023).

A Unevaluated variants

The code base contains two further variants of the model. Neither is evaluated in this paper, and no result in the main text depends on them.

FiLM conditioning. Feature-wise Linear Modulation (FiLM) [31] conditions each GNN layer on an external context. A small generator network maps a global noise level to per-layer affine parameters $(\gamma^{(\ell)}, \beta^{(\ell)})$, which modulate the aggregated messages as $\gamma^{(\ell)} \odot \mathbf{m}^{(\ell)} + \beta^{(\ell)}$ (Fig. 1). In the implementation the generator receives the physical error rate p of the shot, i.e. the true global rate rather than an estimate from the syndrome. This input is one scalar shared by the whole graph and contains no per-qubit reliability information. Proposition 1 shows that uniformly rescaling the channel prior cannot change the hard decisions of unclipped min-sum BP; it does not imply that a GNN cannot use a global scalar as context.

Neural BP. Differentiable BP generalises Eqs. (2)–(3) with learnable per-iteration weights on the channel LLR, the variable-to-check messages and the check-to-variable messages, $w^{(t)} = \mathrm{softplus}(w^{(t)}_{\mathrm{raw}})/\ln 2$. Dividing by $\ln 2$ makes $w_{\mathrm{raw}} = 0$ give $w = 1$ exactly, so at initialisation the unrolled network is standard min-sum BP with scaling factor 0.8.

Algorithm 1: Interleaved GNN-BP Decoding

Input: $\mathbf{s}$, channel LLR Λ, global noise level p

Output: Error estimate $\hat{\mathbf{e}}$

$(\gamma, \beta) \leftarrow \text{FiLMGen}(p)$ `// Conditioning`

$\Delta^{(1)} \leftarrow \textsc{TannerGNN}(\Lambda, \mathbf{s}, \gamma, \beta)$

$\Lambda' \leftarrow \Lambda + \Delta^{(1)}$ `// Prior correction`

$\Lambda_{\text{mid}} \leftarrow \text{NeuralBP}(\Lambda', \mathbf{H}, T_1)$

$p_i^{\text{mid}} \leftarrow \sigma(-\Lambda_{\text{mid},i})\ \forall i$

$\Delta^{(2)} \leftarrow \textsc{TannerGNN}(\Lambda_{\text{mid}}, p^{\text{mid}}, \mathbf{s}, \gamma, \beta)$

$\Lambda'' \leftarrow \Lambda_{\text{mid}} + \Delta^{(2)}$ `// Mid correction`

$\Lambda_{\text{final}} \leftarrow \text{NeuralBP}(\Lambda'', \mathbf{H}, T_2)$

$\hat{e}_i \leftarrow \mathbb{1}[\Lambda_{\text{final},i} < 0]\ \forall i$ **return** $\hat{\mathbf{e}}$

Interleaved pipeline. Algorithm 1 and Fig. 8 give the interleaved pipeline. The GNN intervenes twice: before BP to correct channel priors, and after T_1 iterations using the intermediate posterior information. Qubits with $p_i^{\text{mid}} \approx 0.5$ have low posterior confidence and can receive a second learned correction. Unlike the model of Sec. 3.2, both correction stages are trained jointly through the differentiable BP unroll, so the loss is computed on BP's output marginals. The FiLM and interleaved variants are trained on data regenerated each epoch with $p \sim \text{Uniform}[0.01, 0.10]$, with the focal loss on both components and a syndrome-consistency term, the binary cross-entropy between the measured syndrome and $\frac{1}{2}(1 - \prod_{j \in \mathcal{N}(c)}(1 - 2q_j))$ for each check c, with weight 0.1.

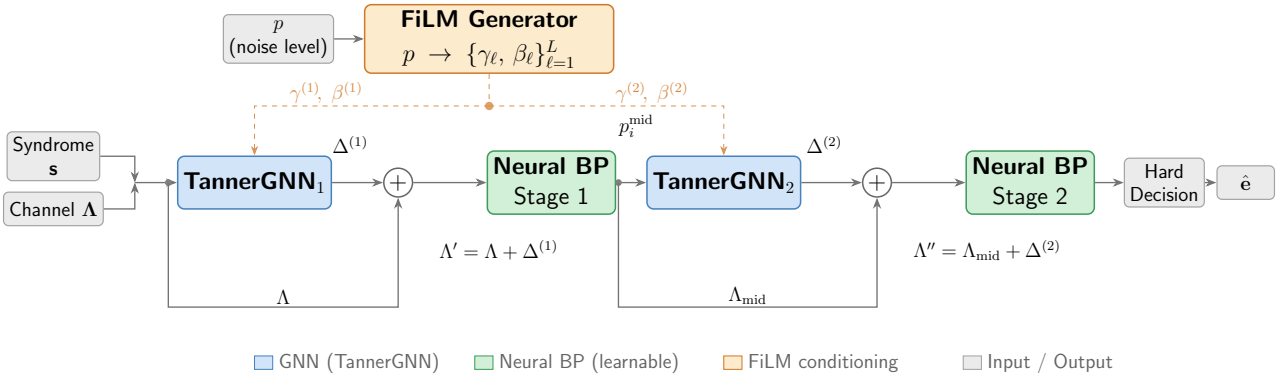


Figure 8: Interleaved GNN-BP pipeline. The GNN applies two successive LLR corrections: one before BP initialisation and one after T_1 iterations, using the mid-decoding posterior as an additional input. Both stages are trained end-to-end through the BP unroll. This variant is not evaluated in this paper.